



香港電腦奧林匹克競賽
Hong Kong Olympiad in Informatics

Data Structures (III)

Bryan Chung {kobebryan9}

2024-03-23

Reference

- This slide is mainly adapted from Data Structures (III) slides (2023) by Gabriel
 - <https://assets.hkoi.org/training2023/ds-iii.pdf>
- OI Wiki
 - <https://oi-wiki.org/>
- Algorithms for Competitive Programming (CP-Algorithms)
 - <https://cp-algorithms.com/index.html>
- ITMO Academy: pilot course in Codeforces
 - <https://codeforces.com/edu/course/2>

Agenda

Sparse Table

- Range Minimum Query (RMQ)
- Binary Lifting

Segment Tree

- Update & Query
- Lazy Propagation

Binary Indexed Tree (Fenwick Tree)

- “Prefix Sum with Efficient Update”

Problem list (Segment Tree)

<https://codeforces.com/blog/entry/22616>

<https://codeforces.com/blog/entry/71925>

<https://codeforces.com/edu/course/2/lesson/4>

<https://codeforces.com/edu/course/2/lesson/5>

<https://codeforces.com/blog/entry/57319>

<https://codeforces.com/contest/438/problem/D>

<https://codeforces.com/gym/104090/problem/M>

<https://loj.ac/p/2269>

Range Minimum Query (RMQ)

Given an array A of N integers and Q queries

Given L and R for each query, find the minimum value in $A[L]$, $A[L + 1]$, ..., $A[R]$

E.g. $A = \{3, 4, 1, 5, 2\}$

$L = 1, R = 2 \rightarrow \text{min value} = 3$

$L = 2, R = 4 \rightarrow \text{min value} = 1$

$L = 4, R = 5 \rightarrow \text{min value} = 2$

Range Minimum Query (RMQ)

Naive Solution: for every query, loop from L to R and take min

Time Complexity: $O(QN)$

Time Limit Exceeded when QN is large :(

How can we do better?

Range Minimum Query (RMQ)

Answer: use sparse table for sure :)

Sparse Table

Let's solve an easier problem first: assume every query's range is 2^x for some x

E.g. $L = 2, R = 9 \rightarrow \text{range} = 9 - 2 + 1 = 8$, which is 2^3

If we precompute all possible ranges for every starting position, we can answer the query in $O(1)$

Formally, we want to precompute $f(i, x) = \min(A[j] \mid i \leq j < (i + 2^x))$

E.g. $f(3, 2) = \min\{A[3], A[4], A[5], A[6]\}$

Sparse Table

Since $x \leq \log_2(N)$, we only need $O(N \log N)$ space to store all possible values

Also, we can compute all values in $O(N \log N)$

Main idea: compute $f(1..N, x + 1)$ from $f(1..N, x)$

Sparse Table

Consider $A = \{3, 4, 1, 5, 2\}$

For $x = 0$,

$f(i, 0) = \min \text{ value from } A[i] \text{ to } A[i + 2^0 - 1] = A[i]$

E.g. $f(2, 0) = A[2]$

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$x = 2$					
$x = 1$					
$x = 0$	3	4	1	5	2

Sparse Table

Consider $A = \{3, 4, 1, 5, 2\}$

Now for $x = 1$,

$f(i, 1) = \text{min value from } A[i] \text{ to } A[i + 2^1 - 1] = \min(A[i], A[i + 1])$

$\therefore f(1, 1) = \min(A[1], A[2])$

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$x = 2$					
$x = 1$	3				
$x = 0$	3	4	1	5	2

Sparse Table

$$f(2, 1) = \min(A[2], A[3])$$

	i = 1	i = 2	i = 3	i = 4	i = 5
x = 2					
x = 1	3	1			
x = 0	3	4	1	5	2

Sparse Table

	i = 1	i = 2	i = 3	i = 4	i = 5
x = 2					
x = 1	3	1	1		
x = 0	3	4	1	5	2

Sparse Table

$$f(4, 1) = \min(A[4], A[5])$$

	i = 1	i = 2	i = 3	i = 4	i = 5
x = 2					
x = 1	3	1	1	2	
x = 0	3	4	1	5	2

Sparse Table

We don't need $f(5, 1) = \min(A[5], A[6])$ as we don't have $A[6]$

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$x = 2$					
$x = 1$	3	1	1	2	
$x = 0$	3	4	1	5	2

Sparse Table

Now for $x = 2$,

$f(i, 2) = \min \text{ value from } A[i] \text{ to } A[i + 2^2 - 1] = \min\{A[i], A[i + 1], A[i + 2], A[i + 3]\}$

Instead of looping from i to $i + 3$, we can compute $f(i, 2)$ from $f(i, 1)$ and $f(i + 2, 1)$!

$f(1, 2) = \min(f(1, 1), f(3, 1))$, $f(1, 1) = \min(A[1], A[2])$, $f(3, 1) = \min(A[3], A[4])$

	$i = 1$	$i = 2$	$i = 3$	$i = 4$	$i = 5$
$x = 2$	1				
$x = 1$	3	1	1	2	
$x = 0$	3	4	1	5	2

Sparse Table

In general, $f(i, x + 1) = \min(f(i, x), f(i + 2^x, x))$

$\therefore f(2, 2) = \min(f(2, 1), f(4, 1))$

	i = 1	i = 2	i = 3	i = 4	i = 5
x = 2	1	1			
x = 1	3	1	1	2	
x = 0	3	4	1	5	2

Sparse Table

We solved the easier version where every query's range is 2^x

How about for queries with arbitrary range?

Sparse Table

Observation: $\min\{a, b, c\} = \min(\min(a, b), \min(b, c))$

Overlapped ranges does not affect the result of min value!

So we just need **two** values from the sparse table:

$\min(A[L..y])$ & $\min(A[x..R])$, and x can be $\leq y$

Sparse Table

E.g. $A = \{3, 4, 1, 5, 2\}$, $L = 2$, $R = 4$

$$f(2, 1) = \min(A[2], A[3])$$

$$f(3, 1) = \min(A[3], A[4])$$

$$\text{Answer} = \min(f(2, 1), f(3, 1))$$

Sparse Table

Let k be the maximum integer such that $R - L + 1 \geq 2^k$

$[L, L + 2^k - 1]$ and $[R - 2^k + 1, R]$ must cover all positions from L to R

$\therefore \text{Answer} = \min(f(L, k), f(R - 2^k + 1, k))$

Sparse Table

E.g. $L = 5, R = 16 \rightarrow R - L + 1 = 12$

$k = 3$ ($2^3 = 8 \leq 12, 2^4 = 16 > 12$)

$f(5, 3) = \min(A[5..12])$

$f(9, 3) = \min(A[9..16])$

Answer = $\min(f(5, 3), f(9, 3))$

Pseudocode

```
precompute()  
  for i = 1 to N  
    ST[i][0] = A[i]  
  for x = 0 to  $\lfloor \log_2(N) \rfloor - 1$   
    for i = 1 to  $N - (2^x + 1 - 1)$   
      ST[i][x + 1] = min(ST[i][x], ST[i +  $2^x$ ][x])
```

Pseudocode

```
query(L, R)
```

```
     $k = \lfloor \log_2(R - L + 1) \rfloor$ 
```

```
    return  $\min(ST[L][k], ST[R - 2^k + 1][k])$ 
```

Reminder: in order to calculate \log_2 in C++ efficiently, either use `std::lg`,
or precompute by $\log N[1] = 0$, $\log N[i] = \log N[i / 2] + 1$

Sparse Table

We solved the problem with $O(N \log N)$ precomputation and $O(1)$ query!

If you still remember, the reason why we can have $O(1)$ query is that **overlapped ranges** does not affect the result of min value

Therefore, as long as the value of an operation would not be affected by **overlapped ranges**, we can have $O(1)$ query using sparse table

- max / and / or / gcd
- any operation that is [idempotent](#)

Sparse Table

Although we cannot have $O(1)$ query for sum / product / xor, we can still have $O(\log N)$ query, as we only need at most $O(\log N)$ values from the sparse table

E.g. $L = 7, R = 19 \rightarrow R - L + 1 = 13 \rightarrow 1101_{(2)}$

We need $[7, 14], [15, 18] \& [19, 19]$ (i.e. $f(7, 3), f(15, 2) \& f(19, 0)$)

Pseudocode

```
query(L, R)
    range = R - L + 1
    sum = 0
    for i =  $\lfloor \log_2(\text{range}) \rfloor$  downto 0
        if  $i^{\text{th}}$  bit of range is 1
            sum += ST[L][i]
            L +=  $2^i$ 
    return sum
```

Binary Lifting

In fact, the way we compute the sum with sparse table is called “binary lifting”

This technique is very useful in many different problems

One of them is lowest common ancestor (LCA)

Lowest Common Ancestor (LCA)

$f(u, i) = 2^i$ -th ancestor of node u

$f(u, 0) = \text{parent of } u$

For finding LCA of u & v , without loss of generality, assume $\text{dep}(v) \geq \text{dep}(u)$

First, we lift node v up such that $\text{dep}(v) = \text{dep}(u)$


If $v = u$, then u is the LCA

Otherwise, for i from high to low, if $f(u, i) \neq f(v, i)$, we lift up node u and node v together,
i.e. $u = f(u, i)$ & $v = f(v, i)$

At last, $f(u, 0)$ is the answer

More details: attend next week's Graph(III) training sessions, or refer to [last year's slide](#)

Sparse Table on Tree



Practice Problems

<https://judge.hkoi.org/task/M2112>

<https://judge.hkoi.org/task/T181>

<https://judge.hkoi.org/task/NP1313>

<https://judge.hkoi.org/task/T114>

Further Readings

https://cp-algorithms.com/data_structures/sparse-table.html

<https://oi-wiki.org/ds/sparse-table/>

<https://oi-wiki.org/topic/rmq/>

Segment Tree

Now we take a look at the range sum query problem with update

Given an array A of N integers and Q operations

Type 1: given x & val , update $A[x] = val$

Type 2: given L & R , query sum $A[L..R]$

Since sparse table doesn't support update, every type 1 operation requires recomputing the table, which is not efficient enough to pass the time limit in usual

Segment Tree

Segment Tree is a more flexible data structure for solving this problem

It is a binary tree

Every node store the information of an interval $[L, R]$


Let $mid = L + (R - L) / 2$

Its left child store the information of $[L, mid]$

Its right child store the information of $[mid + 1, R]$

Segment Tree

$A = \{3, 1, 2, 5, 6, 4, 3, 2\}$



Source: <https://codeforces.com/edu/course/2/lesson/4>

Segment Tree

First, let's walk through how to build the segment tree for an array A

It can be constructed using recursion

In general, if current node's id is x ,

$x * 2$ is used as left child's id

$x * 2 + 1$ is used as right child's id

So we can compute the id directly and don't have to store them separately

Pseudocode

```
build(id, L, R)
    if L = R
        Node[id]  $\leftarrow$  A[L]
        return
    mid  $\leftarrow$  L + (R - L) / 2
    build(id * 2, L, mid)
    build(id * 2 + 1, mid + 1, R)
    Node[id]  $\leftarrow$  Node[id * 2] + Node[id * 2 + 1]

build(1, 1, N)
```

Segment Tree

For length of $A = 2^x$, total number of nodes = $2N - 1$

For length of $A \neq 2^x$, largest id can exceed $2N$, so we usually declare an array of size $4N$ for the segment tree

Time & Space Complexity: $O(N)$

Segment Tree

Now, let's walk through how to query efficiently using segment tree

Again, it can be done by recursion

Pseudocode


```
query(id, L, R, QL, QR) // range [L, R], query range [QL, QR]
    if QR < L or R < QL // no intersection between [L, R] & [QL, QR]
        return 0
    if QL ≤ L and R ≤ QR // [L, R] is fully inside [QL, QR]
        return Node[id]
    mid ← L + (R - L) / 2
    return query(id * 2, L, mid, QL, QR) + query(id * 2 + 1, mid + 1, R, QL,
QR)

query(1, 1, N, QL, QR)
```


Segment Tree

The time complexity for query is $O(\log N)$

Why? Because in each level of the segment tree, we will only use at most two nodes. Notice that for each node we use in each level, they must be consecutive. So if there are three nodes, we can merge two of the three nodes and use their parent instead.



Source: <https://codeforces.com/edu/course/2/lesson/4>

Segment Tree

Finally, let's walk through how to update in the segment tree

Yes, we can still use recursion to do it

Pseudocode

```
update(id, L, R, x, val)
    if L = R
        Node[id]  $\leftarrow$  val
        return
    mid  $\leftarrow$  L + (R - L) / 2
    if x  $\leq$  mid
        update(id * 2, L, mid, x, val)
    else
        update(id * 2 + 1, mid + 1, R, x, val)
    Node[id]  $\leftarrow$  Node[id * 2] + Node[id * 2 + 1]

update(1, 1, N, x, val)
```

Segment Tree

We solved range query with point update (also point query with range update)

For segment tree, the information stored in the node is much more flexible than sparse table, since there is no overlapped intervals

You can store prefix min / max, hash sum, dp table, etc., as long as the operation satisfies some properties of [monoid](#)

Practice Problems

<https://judge.hkoi.org/task/M0921>

<https://judge.hkoi.org/task/M0923>

<https://judge.hkoi.org/task/T152>

<https://codeforces.com/contest/438/problem/D>

Lazy Propagation

Now let's deal with both range update & range query

Given QL, QR, val, update $A[i] += \text{val}$ for $i = \text{QL}$ to QR

We can't simply solve it with the code above since it involves both range query and range update (why?)

Lazy Propagation

We can solve this problem lazily

Instead of updating every index, we store the intermediate information in some intervals, which are exactly the intervals covered in range query

For each node, only when we need to access its children, we propagate the information to the children

That's why it's called “lazy propagation”

Pseudocode

```
push_down(id, L, mid, R)
    Node[id * 2] += lazy[id] * (mid - L + 1)
    Node[id * 2 + 1] += lazy[id] * (R - mid)
    lazy[id * 2] += lazy[id]
    lazy[id * 2 + 1] += lazy[id]
    lazy[id] = 0
```

`lazy[id]` is the value we have updated in current node `id`, but not in its children

Pseudocode

```
query(id, L, R, QL, QR) // range [L, R], query range [QL, QR]
    if QR < L or R < QL // no intersection between [L, R] & [QL, QR]
        return 0
    if QL ≤ L and R ≤ QR // [L, R] is fully inside [QL, QR]
        return Node[id]
    mid ← L + (R - L) / 2
    push_down(id, L, mid, R) // push down only when we need to access children
    return query(id * 2, L, mid, QL, QR) + query(id * 2 + 1, mid + 1, R, QL,
QR)

query(1, 1, N, QL, QR)
```

Pseudocode

```
update(id, L, R, QL, QR, val) // range [L, R], update range [QL, QR]
    if QR < L or R < QL // no intersection between [L, R] & [QL, QR]
        return
    if QL ≤ L and R ≤ QR // [L, R] is fully inside [QL, QR]
        Node[id] += val * (R - L + 1) // update fully covered interval
        lazy[id] += val // store intermediate information in fully covered interval
        return
    mid ← L + (R - L) / 2
    push_down(id, L, mid, R) // push down only when we need to access children
    update(id * 2, L, mid, QL, QR, val)
    update(id * 2 + 1, mid + 1, R, QL, QR, val)
    Node[id] ← Node[id * 2] + Node[id * 2 + 1]

update(1, 1, N, QL, QR, val)
```

Practice Problems

<https://judge.hkoi.org/task/T192>

<https://judge.hkoi.org/task/T213>

<https://codeforces.com/contest/446/problem/C>

Further Readings

https://cp-algorithms.com/data_structures/segment_tree.html

<https://oi-wiki.org/ds/seg/>

<https://oi-wiki.org/geometry/scanning/>

<https://codeforces.com/edu/course/2/lesson/4>

<https://codeforces.com/edu/course/2/lesson/5>

<https://codeforces.com/blog/entry/18051>

<https://www.luogu.com.cn/blog/cyffff/talk-about-segment-trees-split>

<https://www.luogu.com.cn/blog/foreverlasting/xian-duan-shu-fen-zhi-zong-jie>

https://atcoder.github.io/ac-library/production/document_en/segtree.html

https://atcoder.github.io/ac-library/production/document_en/lazysegtree.html

<https://codeforces.com/blog/entry/57319>

Binary Indexed Tree (Fenwick Tree)

All the stuffs you can do with BIT can be done with segment tree

So what the advantages of BIT?

- Less code
- Less space
- Smaller constant factor

Binary Indexed Tree (Fenwick Tree)

As you can see in the agenda, it is described as “Prefix Sum with Efficient Update”

The reason why it is efficient is that it utilizes the binary representation of the id

Binary Indexed Tree (Fenwick Tree)

Let $\text{lowbit}(x)$ be the value of the rightmost bit in binary representation of x

E.g. $x = 22 = 10110_{(2)}$, $\text{lowbit}(x) = 00010_{(2)} = 2$

In BIT, node x stores the information of interval $[x - \text{lowbit}(x) + 1, x]$

How to compute $\text{lowbit}(x)$ efficiently?

Answer: $\text{lowbit}(x) = x \& -x$


Yes, $O(1)$ operation

Binary Indexed Tree (Fenwick Tree)

Let's solve the range sum query problem again with BIT this time

Source:

<https://www.hackerearth.com/practice/notes/binary-indexed-tree-or-fenwick-tree/>



The value in the enclosed box represents $BIT[index]$.

Pseudocode

```
add(id, val)
    while id ≤ N
        Node[id] += val
        id += id & -id
```

```
sum(id)
    res = 0
    while id > 0
        res += Node[id]
        id -= id & -id
    return res
```

Practice Problem

<https://codeforces.com/problemset/problem/830/B>

2D Data Structure

We can extend segment tree / BIT into a 2D data structure, where each node is another segment tree / BIT

Pseudocode

```
add(x, y, val)
  while x ≤ N
    tmp = y
    while y ≤ M
      Node[x][y] += val
      y += y & -y
    x += x & -x
  y = tmp
```

```
sum(x, y)
  res = 0
  while x > 0
    tmp = y
    while y > 0
      res += Node[x][y]
      y -= y & -y
    x -= x & -x
  y = tmp
  return res
```

2D Data Structure

2D data structure generally has a higher time complexity and memory storage

It is quite rare to see a problem that requires 2D data structure to solve

But still, if you know how to implement 1D data structure, 2D data structure is not that difficult to implement, although sometime it is quite tedious

Practice Problem

<https://judge.hkoi.org/task/I0111>

Further Readings

https://cp-algorithms.com/data_structures/fenwick.html

<https://oi-wiki.org/ds/fenwick/>

<https://www.luogu.com.cn/blog/kingxbz/shu-zhuang-shuo-zu-zong-ru-men-dao-ru-fen>

<https://www.luogu.com.cn/blog/countercurrent-time/qian-tan-shu-zhuang-shuo-zu-you-hua>